

# **CS / EE 320**

# **Computer Organization and Assembly Language**

## **Lecture 25**

## **Spring 2024**

**Shahid Masud**

**Topics: Cache/Memory Update Strategies, Cache  
Performance, Multi Level Caches**

# Topics

- Examples of calculating different parameters in Cache Design
- Cache Replacement: LRU, LFU, FIFO, Random schemes
- Cache Writes – Write Back and Write Through schemes
- Some Examples of Cache Performance
- Introduction to Multi-level Caches
- Examples of Cache Performance

**Quiz 6 Next Monday**

# Dealing with Cache Miss and Cache Line Replacement

# Cache Misses

- On cache hit, CPU proceeds normally
- On cache miss
  - Stall the CPU pipeline
  - Fetch block from next level of hierarchy
  - Instruction cache miss
    - Restart instruction fetch
  - Data cache miss
    - Complete data access

# Sources of Misses

- Compulsory misses (aka cold start misses)
  - First access to a block
- Capacity misses
  - Due to finite cache size
  - A replaced block is later accessed again
- Conflict misses (aka collision misses)
  - In a non-fully associative cache
  - Due to competition for entries in a set
  - Would not occur in a fully associative cache of the same total size

# Replacement Algorithms (1) Direct mapping

- No choice
- Each block only maps to one line
- Replace that line

# Replacement Algorithms (2)

## Associative & Set Associative



- Hardware implemented algorithm (speed)
- Least Recently used (LRU)
- e.g. in 2 way set associative
  - Which of the 2 block is lru?
- First in first out (FIFO)
  - replace block that has been in cache longest
- Least frequently used
  - replace block which has had fewest hits
- Random

# Write Policy

- Must not overwrite a cache block unless main memory is up to date
- Multiple CPUs may have individual caches
- I/O may address main memory directly



# Write-through scheme

- All writes go to main memory as well as cache
- Multiple CPUs can monitor main memory traffic to keep local (to CPU) cache up to date
- Lots of traffic
- Slows down writes

# Write-Through

- On data-write hit, could just update the block in cache
  - But then cache and memory would be inconsistent
- Write through: also update memory
- But makes writes take longer
  - e.g., if base CPI = 1, 10% of instructions are stores, write to memory takes 100 cycles
    - Effective CPI =  $1 + 0.1 \times 100 = 11$
- Solution: write buffer
  - Holds data waiting to be written to memory
  - CPU continues immediately
    - Only stalls on write if write buffer is already full

# Write-back

- Updates initially made in cache only
- Update bit for cache slot is set when update occurs
- If block is to be replaced, write to main memory only if update bit is set
- Other caches get out of sync
- I/O must access main memory through cache
- N.B. 15% of memory references are writes

# Write-Back

- Alternative: On data-write hit, just update the block in cache
  - Keep track of whether each block is dirty
- When a dirty block is replaced
  - Write it back to memory
  - Can use a write buffer to allow replacing block to be read first

# Write Miss

- What should happen on a write miss?
- Alternatives for write-through
  - Allocate on miss: fetch the block
  - Write around: don't fetch the block
    - Since programs often write a whole block before reading it (e.g., initialization)
- For write-back
  - Usually fetch the block

# Block Placement

- Determined by associativity
  - Direct mapped (1-way associative)
    - One choice for placement
  - n-way set associative
    - n choices within a set
  - Fully associative
    - Any location
- Higher associativity reduces miss rate
  - Increases complexity, cost, and access time


# Cache Design Trade-offs



| Design change          | Effect on miss rate        | Negative performance effect   |
|------------------------|----------------------------|---|
| Increase cache size    | Decrease capacity misses   | May increase access time  |
| Increase associativity | Decrease conflict misses   | May increase access time  |
| Increase block size    | Decrease compulsory misses | Increases miss penalty. For very large block size, may increase miss rate due to pollution. |

# Cache Control

- Example cache characteristics
  - Direct-mapped, write-back, write allocate
  - Block size: 4 words (16 bytes)
  - Cache size: 16 KB (1024 blocks)
  - 32-bit byte addresses
  - Valid bit and dirty bit per block
  - Blocking cache
    - CPU waits until access is complete





# Multilevel On-Chip Caches



| Characteristic         | ARM Cortex-A53   | Intel Core i7                               |
|------------------------|--|---|
| L1 cache organization  | Split instruction and data caches                                    | Split instruction and data caches           |
| L1 cache size          | Configurable 16 to 64 KiB each for instructions/data                 | 32 KiB each for instructions/data per core  |
| L1 cache associativity | Two-way (I), four-way (D) set associative                            | Four-way (I), eight-way (D) set associative |
| L1 replacement         | Random   | Approximated LRU                            |
| L1 block size          | 64 bytes   | 64 bytes                                    |
| L1 write policy        | Write-back, variable allocation policies (default is Write-allocate) | Write-back, No-write-allocate               |
| L1 hit time (load-use) | Two clock cycles   | Four clock cycles, pipelined                |
| L2 cache organization  | Unified (instruction and data)                                       | Unified (instruction and data) per core     |
| L2 cache size          | 128 KiB to 2 MiB   | 256 KiB (0.25 MiB)                          |
| L2 cache associativity | 16-way set associative   | 8-way set associative                       |
| L2 replacement         | Approximated LRU   | Approximated LRU                            |
| L2 block size          | 64 bytes   | 64 bytes                                    |
| L2 write policy        | Write-back, Write-allocate   | Write-back, Write-allocate                  |
| L2 hit time            | 12 clock cycles  | 10 clock cycles                             |
| L3 cache organization  | –  | Unified (instruction and data)              |
| L3 cache size          | –  | 8 MiB, shared                               |
| L3 cache associativity | –  | 16-way set associative                      |
| L3 replacement         | –  | Approximated LRU                            |
| L3 block size          | –  | 64 bytes                                    |
| L3 write policy        | –  | Write-back, Write-allocate                  |
| L3 hit time            | –  | 35 clock cycles                             |

# Concluding Remarks

- Fast memories are small, large memories are slow
  - We really want fast, large memories
  - Caching gives this illusion
- Principle of locality
  - Programs use a small part of their memory space frequently
- Memory hierarchy
  - L1 cache  $\leftrightarrow$  L2 cache  $\leftrightarrow$  ...  $\leftrightarrow$  DRAM memory  $\leftrightarrow$  disk
- Memory system design is critical for multiprocessors

- **Multi Level Caches**
- **Performance of Cache and Memory System**
- **Examples of Cache Performance**

# Example 1 of Cache Performance

## Question

Given that:

**T1** is access time from Cache = 0.01  $\mu$ sec

**T2** is access time from Main Memory = 0.1  $\mu$ sec

95% accesses are from Cache

5% accesses are from Main Memory

Find the Average Time seen by CPU to Access one word?

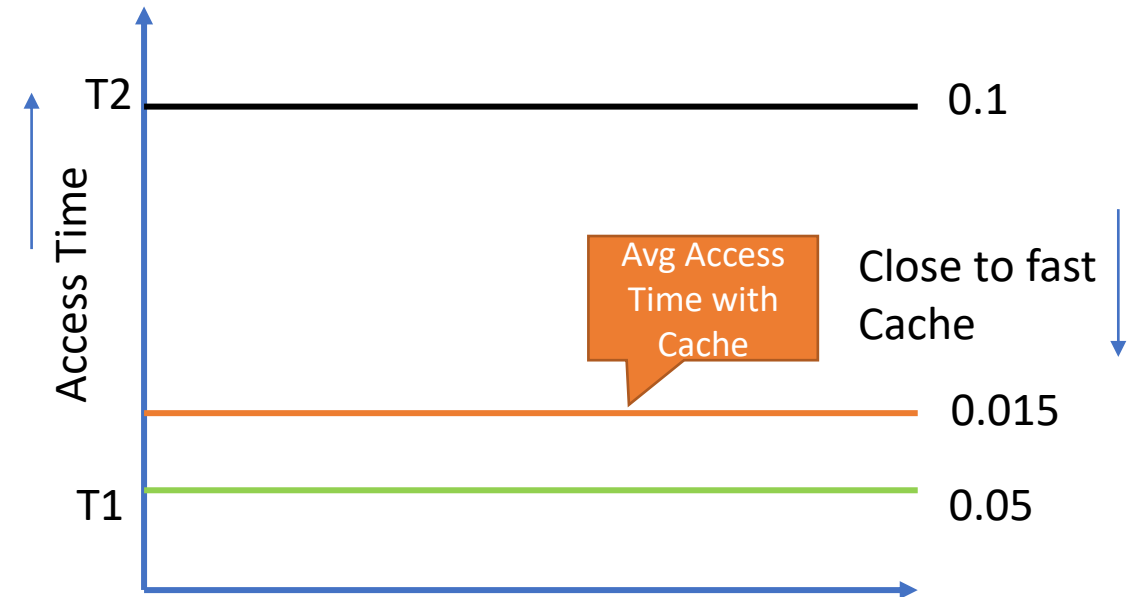
## Solution

Average time to access one word:

$$T_{avg} = (0.95 \times 0.01) + (0.05 \times 0.1)$$


$$T_{avg} = 0.0095 + 0.0055$$

$$T_{avg} = 0.015 \mu\text{sec}$$



# Split Caches


Separate Cache for Instructions and Data to take advantage of pipelining  
In case the program has too many localized accesses to data only or instructions only,  
Then there may be no performance improvement



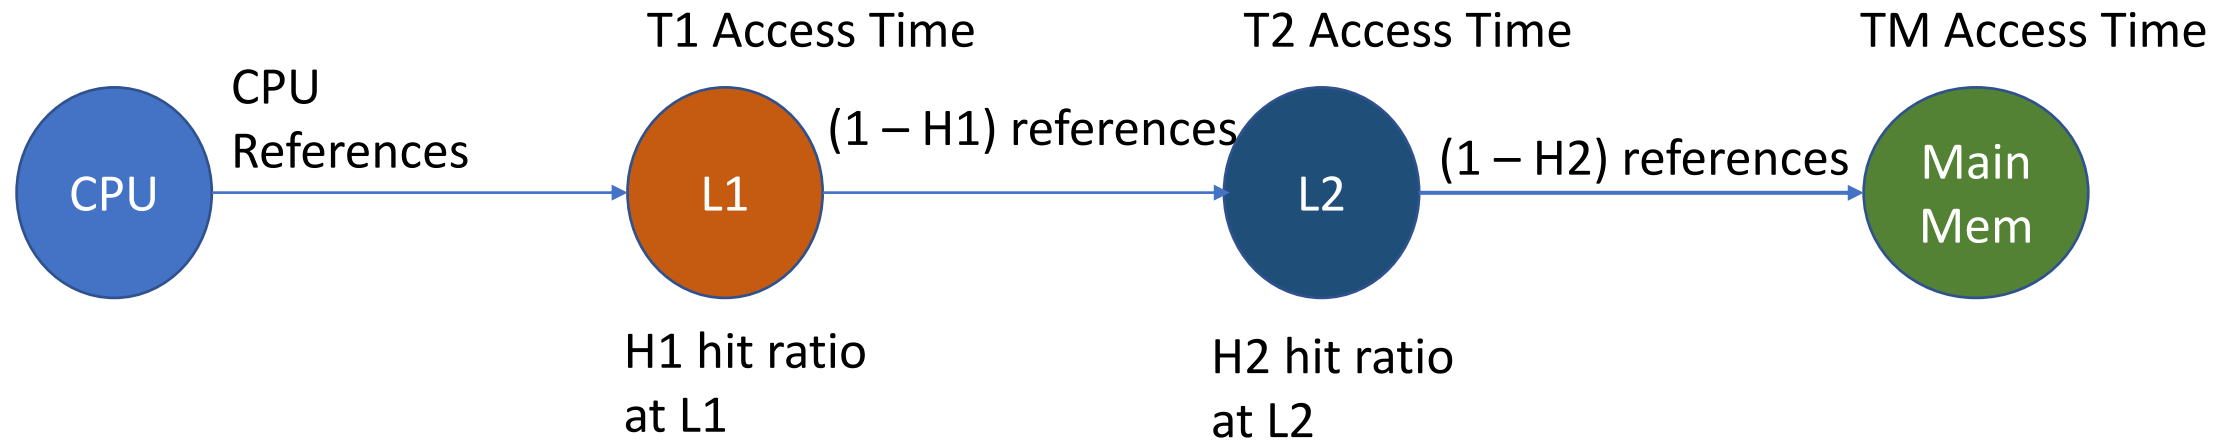
- Primary cache attached to CPU
  - Small, but fast
- Level-2 cache services misses from primary cache
  - Larger, slower, but still faster than main memory
- Main memory services L-2 cache misses
- Some high-end systems include L-3 cache

# Multi Level Caches

Main purpose is to reduce Miss Penalty



# Multi Level Caches - Performance



**Average Access Time** in Multilevel Cache:

$$T_{avg} = T1 + (1 - H1) \times T2 + (1 - H1) \times (1 - H2) \times TM$$

$$(1 - H1) = \text{miss rate at L1} = \text{MRL1}$$

$$T_{avg} = T1 + \text{MRL1} \times T2 + \text{MRL1} \times \text{MRL2} \times TM$$

$$\text{MRL1} \times \text{MRL2} = \text{Global Miss Rate at L2}$$



# Example 1, find Tavg in 2 level cache

- CPU generates 100 memory references. 80 are hit in L1, 20 are miss in L1. 16 are hit in L2 and 4 are miss in L2. Access time of L1 is 1 cycle. Access time of L2 is 10 cycles. Access time of Main Mem is 100 cycle. Find average memory access time Tavg?

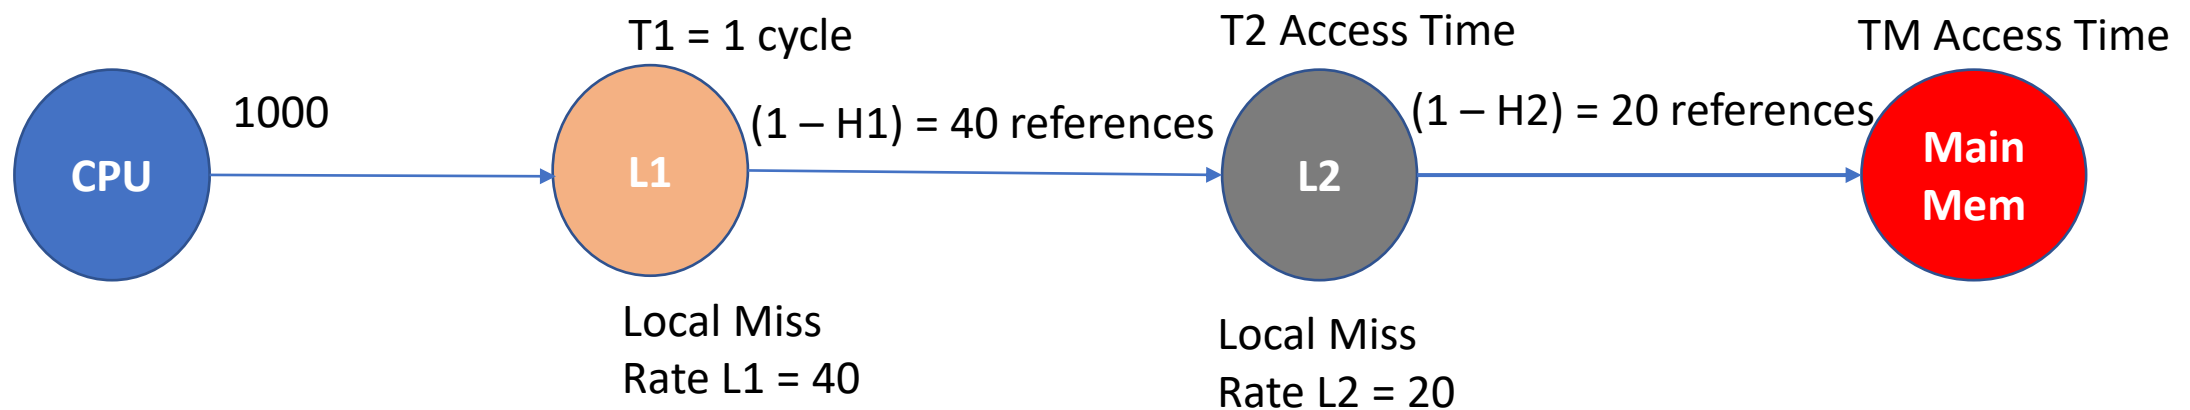
- $T_{avg} = T_1 + MRL_1 \times T_2 + MRL_1 \times MRL_2 \times T_M$
- $= 1 + (20/100) \times 10 + [(20/100) \times (4/20)] \times 100$
- $= 1 + 2 + 4 = 6$  cycles Avg Access time

Note:

MRL1 = Miss Rate at Cache L1

# Example 2, find Tavg in 2 level cache

- CPU generates 1000 memory references. Number of Miss in L1 is 40. Number of miss in L2 is 20. Access time of L1 is 1 cycle. Access time of L2 is 10 cycle. Find Tavg when Miss Penalty of L2 is 100 cycle.



$$\begin{aligned} T_{avg} &= T1 + MRL1 \times T2 + MRL1 \times MRL2 \times TM \\ &= 1 + (40/1000) \times 10 + [(40/1000) \times (20/40)] \times 100 \\ &= 1 + 0.4 + 2 = 3.4 \text{ cycles Avg Access time} \end{aligned}$$

- Components of CPU time
  - Program execution cycles
    - Includes cache hit time
  - Memory stall cycles
    - Mainly from cache misses
- With simplifying assumptions:

Memory stall cycles

$$= \frac{\text{Memory accesses}}{\text{Program}} \times \text{Miss rate} \times \text{Miss penalty}$$

$$= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Misses}}{\text{Instruction}} \times \text{Miss penalty}$$

# Cache Performance Example

- Given
  - I-cache miss rate = 2%
  - D-cache miss rate = 4%
  - Miss penalty = 100 cycles
  - Base CPI (ideal cache) = 2
  - Load & stores are 36% of instructions
- Miss cycles per instruction
  - I-cache:  $0.02 \times 100 = 2$
  - D-cache:  $0.36 \times 0.04 \times 100 = 1.44$
- Actual CPI =  $2 + 2 + 1.44 = 5.44$ 
  - Ideal CPU is  $5.44/2 = 2.72$  times faster

- Hit time is also important for performance
- Average memory access time (AMAT)
  - $AMAT = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$
- Example
  - CPU with 1ns clock, hit time = 1 cycle, miss penalty = 20 cycles, l-cache miss rate = 5%
  - $AMAT = 1 + 0.05 \times 20 = 2\text{ns}$ 
    - 2 cycles per instruction

# Performance Summary

- When CPU performance increased
  - Miss penalty becomes more significant
- Decreasing base CPI
  - Greater proportion of time spent on memory stalls
- Increasing clock rate
  - Memory stalls account for more CPU cycles
- Can't neglect cache behavior when evaluating system performance

# Multilevel Cache CPI Example

- Given
  - CPU base CPI = 1, clock rate = 4GHz
  - Miss rate/instruction = 2%
  - Main memory access time = 100ns
- With just primary cache
  - Miss penalty =  $100\text{ns} / 0.25\text{ns} = 400$  cycles
  - Effective CPI =  $1 + 0.02 \times 400 = 9$

# Readings

- Chap 5 of P&H Textbook